

# Multi-GPU training with TensorFlow on Piz Daint

Input pipelines with TensorFlow's `tf.data` API


Rafael Sarmiento and Henrique Mendonça

ETH Zürich / CSCS


7<sup>th</sup>-8<sup>th</sup> September 2020

# Outline

- Input pipelines
- [lab] Building input pipelines
- Read and write TFRecord files
- [lab] Reading and writing TFRecord files




# Pipelining



# Pipelining



# Pipelining




# Input pipelines



## Extract

Read data from persistent storage (HDD, SSD, GCS, HDFS, ...)



## Transform

Use **CPU cores** to parse and perform preprocessing operations on the data, shuffling and batching



## Load

Load the transformed data onto the accelerator devices that execute the model

# TensorFlow's `tf.data` API

```
# Extract  
dataset = tf.data.TFRecordDataset("./train.tfrecords")
```



# TensorFlow's `tf.data` API

# Extract

```
dataset = tf.data.TFRecordDataset("./train.tfrecords")
```

# Transform

```
dataset = dataset.shuffle(1000)
```

```
dataset = dataset.batch(64)
```

```
dataset = dataset.map(resize_images)
```

```
dataset = dataset.repeat(100)
```

# TensorFlow's `tf.data` API

# Extract

```
dataset = tf.data.TFRecordDataset("./train.tfrecords")
```

# Transform

```
dataset = dataset.shuffle(1000)
```

```
dataset = dataset.batch(64)
```

```
dataset = dataset.map(resize_images)
```

```
dataset = dataset.repeat(100)
```

# 'Load' happens automatically!

# [lab] Building input pipelines

Let's open an empty notebook and create some simple input pipelines. A good starting point can be to generate random numpy data and create the pipeline including maps, filters, batching, shuffling and repeating operations. `tf.data.Dataset.from_tensor_slices` can be used for the *extract* phase.

Let's start then with the notebook `1_getting_started_with_tensorflows_dataset_api.ipynb`, that's on the folder `input_pipelines/`.

Then we continue to the rest of the notebooks `*_getting_started_with_tensorflows_dataset_api.ipynb` to try other examples.

# TFRecord files

- TFRecord is the format of data storage recommended for TensorFlow
- TFRecord is a simple record-oriented binary format
- Data is serialized with protocol buffers and stored as collections of meaningful units (records) in contrast to a byte-oriented filesystem, where the data is treated as an unformatted stream of bytes
- On the deep learning context each record would be an item of the dataset

# Best practices for storing data

- Avoid storing data as multiple small files.
- Divide large datasets ( $\gg 1$  Gb) into TFRecord files of about 150 Mb.
- For smaller data sets (200MB-1GB), a single TFRecord is enough.
- Randomly divide the data into multiple files and then shuffle the filenames uniformly. This helps to decrease the size of the shuffle buffer.
- Read the data from `$SCRATCH`.


# [lab] Reading and writing TFRecord files

Let's run the notebook `read_and_write_TFRecord_files.ipynb`. We are going to write two cat images to a TFRecord file and then we are going to read them to check that they can be recovered correctly.

Following the same steps, write the MNIST dataset to a TFRecord file. We will use it later on a lab.

You can get the MNIST data as numpy arrays with

```
from tensorflow.keras.datasets import mnist
(x_train, y_train), (x_test, y_test) = mnist.load_data()
```



# Thank you for your attention!